

• METHODOLOGY V2

claude code · self-feeding delivery · july 2026

survey-corps → **war-machine** →
war-campaign → **aftermath**

spec → plan → red-team → war

v1 was a pipeline you drive. v2 is a machine that feeds itself – you ratify decisions and review PRs.

• survey-corps

• war-machine

• war-campaign

• aftermath

• lessons-learned

WORK · AUDIT · REFINE

V1 RECAP · JUNE 2026

The inner loop still runs



This still runs. It's now the inner loop.

Strict roles, unchanged: grill authors · war-strategy converts; red-team validates and never converts · war executes.

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war


What changed in a month

V1 · JUNE 2026		V2 · JULY 2026
commands	4	11
specs	you, grilled down every branch	/survey-corps mines memories + issues
plans	you, grilled again	the machine grills itself
execution	one plan · one /war	campaigns: N plans, unattended
scheduling	a hand-written overnight /loop	every outer step is cron-able
cleanup	—	/aftermath, evidence-gated
reach	one repo	a four-repo theater

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

THE OUTER LOOP

Five commands, five artifacts



Every step has an autonomous mode. The sequence is cron-able end to end – the inner loop runs inside war-campaign.

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

Where specs come from

01

/survey-corps

docs/specs/ · survey manifest

--erwin pause for group approval

- ▶ **Memories and open issues go in**

It mines the construct and the backlog together, and groups related issues into coherent work.

- ▶ **Grouped, war-shaped specs come out**

Each group lands as a spec the inner pipeline can fight, plus a survey manifest recording the sweep.

- ▶ **Autonomous by default**

--erwin holds the sweep for grouping approval. Without it, the corps surveys on its own.

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

The machine grills itself

02

/war-machine

docs/plans/ • docs/roadmaps/

drafter ⇌ adversarial grill → plan

- ▶ **A drafter + adversarial pair, per spec**

The grill that interviewed you now interrogates the drafter.
Converts specs to plans – never launches, never red-teams.

- ▶ **Only key decisions float up**

You ratify a handful of design choices – or none: --afk writes a provenance-marked AI-Commander's Intent (ADR 0014).

- ▶ **Intent is the ceiling**

Every plan opens with Commander's Intent – Purpose, Method, a numbered checkable End state. The plan slice is the floor, intent is the ceiling.

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

Stack and pLow

03

`/war-campaign <roadmap>`

`.claude/campaigns/ · stacked PRs`

`halt-and-hold` on any failure



Each plan branches from the prior plan's tip (ADR 0011). PRs stack; merge runs bottom-up.

Red-team + war run per plan, back to back, unattended. The campaign never auto-invokes – the operator schedules it.

survey-corps → war-machine → `war-campaign` → aftermath
spec → plan → red-team → war

Surviving the night

●●● .claude/campaigns/.../CAMPAIGN-STATE.md

```
# rewritten before every long wait  
roadmap: 5 plans  
done: plan-1 ✓ · plan-2 ✓  
current: plan-3 · phase 2 · mid-war  
resume: verify against git, not this file
```

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

► Write-ahead, always

The checkpoint rewrites itself before every long wait, so a compaction mid-wait loses nothing.

► Re-injected after compaction

A SessionStart hook re-injects campaign state into the fresh context (ADR 0016). The campaign survives its own amnesia.

► git is the resume source of truth

The checkpoint is a hint. What actually landed is read from the repo – never from ledger claims.

INSIDE /WAR

Work • Audit • Refine

```
/war <plan> --working dev/<slug> --landing master --afk
```

frozen phase base → provision → **work + audit ×waves** → refine → gate-audit → land → **record**

01

WORK

Fresh workers implement each task in isolated git worktrees, writing the mapped tests – wave by wave.

02

AUDIT

1-5 read-only seats judge one pinned SHA through distinct lenses. Critical / Major findings block.

03

REFINE

A serial merge queue rebases, re-runs the gate, and lands approved tasks – one at a time.

04

RECORD

A write-scoped servitor captures the phase's durable learnings to memory, plus a machine-readable handoff.

Repeats per phase. Each phase lands as one `--no-ff commit`; one PR opens at the end. Land is push-first compare-and-swap – never force-push.

Provision is refiner-owned: if it fails, the phase is env-blocked and no worker ever spawns.

survey-corps → war-machine → **war-campaign** → aftermath
spec → plan → red-team → war

QUALITY

The audit gate



HOLDS UNWATCHED This month red-team caught the plan's own false-pass grep — and a git fetch refmap misconception — before execution. (2026-07)

survey-corps → war-machine → **war-campaign** → aftermath
spec → plan → red-team → war

THE CAST

Five roles in the workflow

Lead

`your chat`

Decomposes the plan, proposes, approves, and checks in at each phase.

↳ Mayor

Worker

`opus · max`

Implements one task in an isolated git worktree, writing the mapped tests.

↳ Połecat

Auditor

`opus · xhigh · R0`

Bash locked to an allowlisted git verb set. Severity-tags findings; must be unanimous.

↳ Nun

Refinery

`sonnet · merge queue`

Rebases, re-runs the gate, and lands approved tasks — one at a time.

↳ Refinery

Servitor

`sonnet · write-scoped`

No Bash at all. Records durable learnings; every fact must carry provenance.

↳ bd remember

`red-team: opus`

`a fable fix-agent seat lives in the config`

the Auditor is the "Nun" Gas Town never shipped — WAR built it first-class · config as of 2026-07-17

survey-corps → war-machine → `war-campaign` → aftermath
spec → plan → red-team → war

One fact per file



"Your fiftieth run knows what your fifth one learned."

91 hot + 79 archived committed lessons · as of 2026-07-17

MEMORY.md — generated projection, hard byte budget

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

The provenance ladder



the ladder drives recall ranking and eviction order

TEMPERATURE IS LOCATION

hot = root · cold = archive/. Archiving is a move, never a deletion.

GRADUATION CANDIDATES

A lesson that recurs $\geq 2\times$ with a machine-checkable invariant is nominated for promotion – from memory into an enforced guard.

ONE FACT PER FILE

Every lesson is one Markdown file with provenance frontmatter. No blobs, no drift.

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

Delete only what's proven



Worktrees, task branches, closed-out issues – each deletion carries the evidence that proved it safe.

83

unswept task worktrees across three repos
awaited /aftermath on 2026-07-17.

Cleanup earned its command.

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

GOVERNANCE

The pipeline is gospel

36

ADRs govern the machine · as of 2026-07-17

Drift-guard tests – the doctrine is executable; deviation fails CI.

Version-slot lock – version literals rot; write 0.14.x, never a patch number.

CONTEXT.md – one ubiquitous language across every seat.

THE PAPER TRAIL · WORKAUDITREFINE ·
2026-07-17

- 96 specs
- 97 plans
- 12 roadmaps
- 78 red-team reports
- 64 run ledgers
- 13 run telemetry manifests
- 8 survey manifests
- 91+79 lessons, hot + archived

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

NEW HIGH SCORE · 2026-07-23

Nine plans in. Nine PRs out.

run-resilience-and-hardening · one unbroken campaign session · PRs #1008 → #1077, merge bottom-up

9=9

plans → stacked PRs

35

workflows
dispatched

490

subagents
orchestrated

41.8M 9.4K

tokens consumed

tool calls

✓ zero tasks escalated

✓ zero conflicts

✓ zero holds

9× red-team CLEARED

8 version bumps · 0.14.48 → 0.14.56

A clean sweep — every plan carried from raw spec through red-team, execution, and a stacked PR.

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

THE SCHEDULED CAMPAIGN • 2026-07-16 → 17

One command. Zero interventions.

operator-scheduled /war-campaign over a 5-plan roadmap • checkpoint rewrote itself before every long wait

plan-1 ✓ landed PR #942

plan-2 ✓ landed PR #949

stacked on plan-1

plan-3 ✓ landed PR #956

three releases on the 0.14.x line

plan-4 ● mid-war
at dawn red-teamed itself CLEARED in 3 rounds • phase 1 landed 10/10 End state •
phase 2 releasing

plan-5 ○ queued

SAME WEEK 8 open issues → /survey-corps → 5 specs → /war-machine → 5 plans + roadmap, under 4h of wall clock
(survey 22:06 UTC • plans consumed 01:13 UTC)

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

The machine ported itself

WorkAuditRefine

the home repo

Forges itself: the methodology's own campaigns run here, on the repo that defines them.

96 specs · 97 plans

SQL / AutoIndex

campaign theater

Campaigns ran here too – same doctrine, different codebase, no adaptation layer.

sibling repo

SQL / Otto

war-engine-port

Ran a war-engine-port campaign: the machine used itself to install itself.

sibling repo

ACROSS 4 REPOS WITH WAR STATE · 2026-07-17 102 specs 101 plans 93 red-team reports 4 campaigns

13 campaign plans landed

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war

THE CRONTAB

Run it tonight

/survey-corps

writes the specs

memories + issues, grouped

→ /war-machine

writes the plans

grills itself; you ratify

→ /war-campaign

fights it

red-team + war · stacked
PRs

→ /aftermath

sweeps the field

evidence-gated cleanup

●●● crontab · every step is cron-able

```
# one theater, four lines – the times are yours
```

```
22:00 /survey-corps
```

```
23:00 /war-machine
```

```
00:00 /war-campaign <roadmap>
```

```
07:00 /aftermath
```

You ratify decisions. You review PRs.
It does the rest while you sleep.

The audit gate plus your CI make it safe to run nightly.

survey-corps → war-machine → war-campaign → aftermath
spec → plan → red-team → war